

Regression in R

create some fake data

```
x1 <- 1:100

x2 <- -0.1*x1 + rnorm(100)

x3 <- 0.05*x2 + rnorm(100)

y <- 2*x1 + 10*rnorm(100) + 10*x2

dat <- data.frame( y, x1, x2, x3 )

head( dat )
```

| ## | | y | x1 | x2 | x3 |
|------|--|------------|----|-------------|------------|
| ## 1 | | -1.356088 | 1 | -0.01531433 | -0.9257205 |
| ## 2 | | -0.356096 | 2 | -1.45542891 | 2.0086764 |
| ## 3 | | -8.805591 | 3 | -1.58602538 | 0.1715653 |
| ## 4 | | -24.880864 | 4 | -2.48621082 | -0.8030154 |
| ## 5 | | -3.444461 | 5 | -0.28043740 | -2.2281208 |
| ## 6 | | 8.913825 | 6 | -0.81350619 | 0.4789454 |

descriptive statistics

```
library( pastecs )

## Loading required package: boot

print( t( stat.desc( dat ) ), digits=3 )
```

| ## | nbr.val | nbr.null | nbr.na | min | max | range | sum | median | mean |
|-------|---------|----------|--------|--------|---------|--------|--------|--------|--------|
| ## y | 100 | 0 | 0 | -24.88 | 118.193 | 143.07 | 4906.4 | 52.744 | 49.064 |
| ## x1 | 100 | 0 | 0 | 1.00 | 100.000 | 99.00 | 5050.0 | 50.500 | 50.500 |
| ## x2 | 100 | 0 | 0 | -11.09 | 0.575 | 11.67 | -512.3 | -5.185 | -5.123 |
| ## x3 | 100 | 0 | 0 | -2.23 | 2.446 | 4.67 | -25.9 | -0.328 | -0.259 |

```
##      SE.mean CI.mean.0.95      var std.dev coef.var
## y      3.216      6.381 1034.29   32.16   0.655
## x1      2.901      5.757  841.67   29.01   0.574
## x2      0.299      0.593    8.93    2.99  -0.583
## x3      0.103      0.204    1.06    1.03  -3.971
```

```
# To copy and paste into Excel:
#
# descriptives <- t( stat.desc(dat) )
#
# write.table( descriptives, "clipboard", sep="\t", row.names=TRUE )

# To create nicely formatted tables for markdown documents use the kable() function
```

```
library( knitr )
```

```
kable( t( stat.desc( dat )[ c(1,4,5,8,9,13), ] ), format="markdown", digits=3 )
```

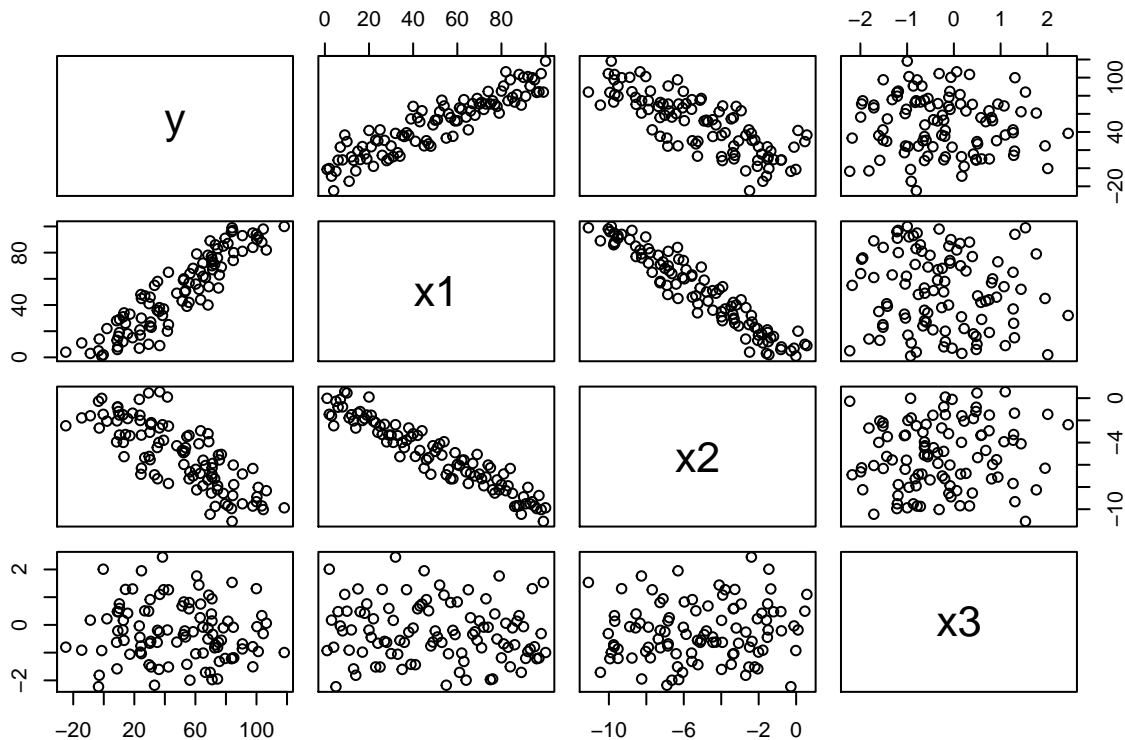
| | nbr.val | min | max | median | mean | std.dev |
|----|---------|---------|---------|--------|--------|---------|
| y | 100 | -24.881 | 118.193 | 52.744 | 49.064 | 32.160 |
| x1 | 100 | 1.000 | 100.000 | 50.500 | 50.500 | 29.011 |
| x2 | 100 | -11.094 | 0.575 | -5.185 | -5.123 | 2.989 |
| x3 | 100 | -2.228 | 2.446 | -0.328 | -0.259 | 1.027 |

| | nbr.val | min | max | median | mean | std.dev |
|----|---------|---------|---------|--------|--------|---------|
| y | 100 | -11.210 | 118.497 | 50.591 | 52.395 | 31.342 |
| x1 | 100 | 1.000 | 100.000 | 50.500 | 50.500 | 29.011 |
| x2 | 100 | -10.851 | 1.071 | -5.097 | -4.964 | 3.135 |
| x3 | 100 | -2.548 | 1.857 | -0.409 | -0.380 | 1.060 |

pretty pairs plot

Convenient visual descriptives:

```
pairs( dat )
```



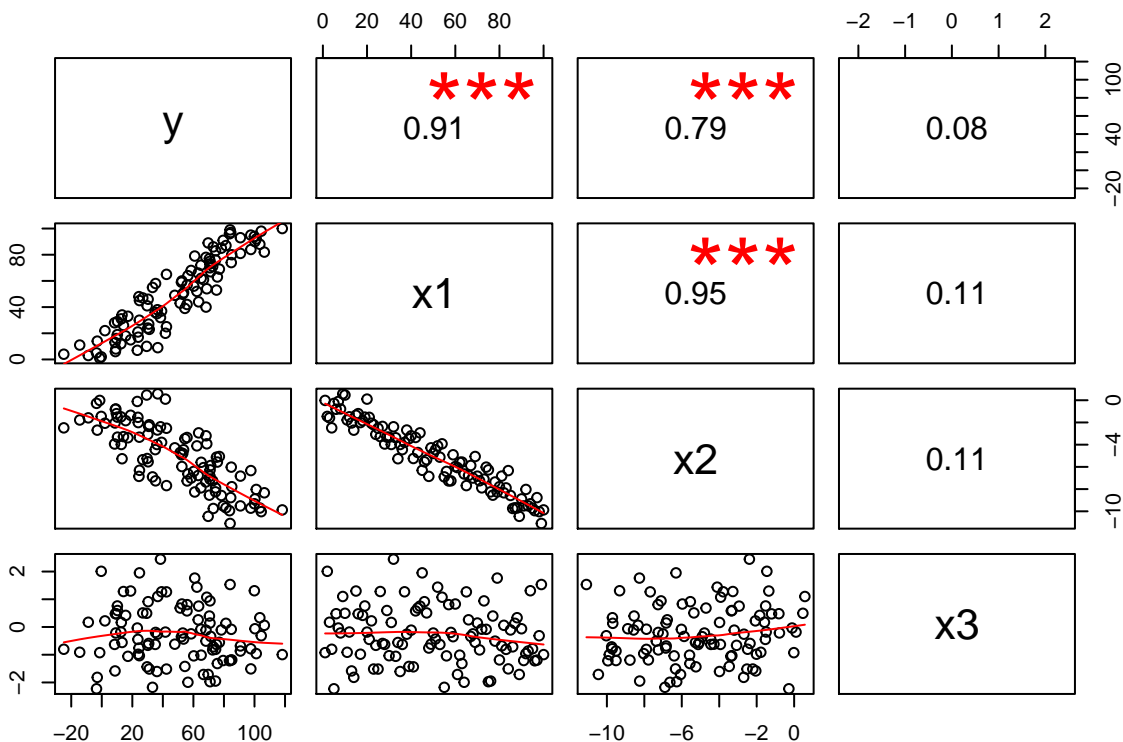
This one is better:

```
panel.cor <- function(x, y, digits=2, prefix="", cex.cor)
{
  usr <- par("usr"); on.exit(par(usr))
  par(usr = c(0, 1, 0, 1))
  r <- abs(cor(x, y))
  txt <- format(c(r, 0.123456789), digits=digits)[1]
  txt <- paste(prefix, txt, sep="")
  if(missing(cex.cor)) cex <- 0.8/strwidth(txt)

  test <- cor.test(x,y)
  # borrowed from printCoefmat
  Signif <- symnum(test$p.value, corr = FALSE, na = FALSE,
    cutpoints = c(0, 0.001, 0.01, 0.05, 0.1, 1),
    symbols = c("***", "**", "*", ".", " "))

  text(0.5, 0.5, txt, cex = 1.5 )
  text(.7, .8, Signif, cex=cex, col=2)
}


pairs( dat, lower.panel=panel.smooth, upper.panel=panel.cor)
```



create some fake regression data

```
x <- 1:100
y <- 2*x + rnorm(100,0,20)

plot( x, y )
```



```
dum <- sample( c("NJ","NY","MA","PA"), 100, replace=T )
```

basic regression syntax

```
lm( y ~ x )
```

```
##
## Call:
## lm(formula = y ~ x)
##
## Coefficients:
## (Intercept)          x
##      -4.460       2.114
```

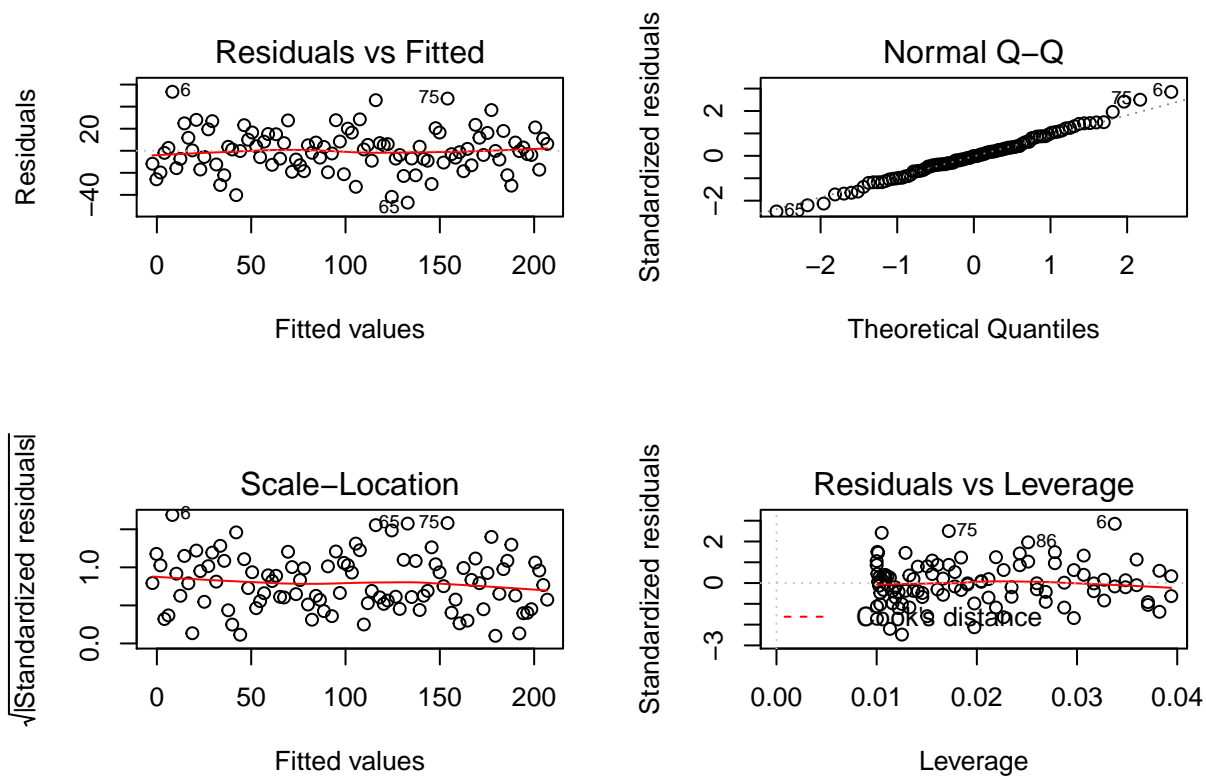
```
m.01 <- lm( y ~ x )
```

```
summary( m.01 )
```

```
##
## Call:
## lm(formula = y ~ x)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -47.000 -11.932  -0.287   11.197   53.447
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -4.46027    3.84369  -1.16   0.249
## x             2.11414    0.06608   31.99 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 19.07 on 98 degrees of freedom
## Multiple R-squared:  0.9126, Adjusted R-squared:  0.9117
## F-statistic: 1024 on 1 and 98 DF,  p-value: < 2.2e-16
```

nice visual diagnostics

```
par( mfrow=c(2,2) )
plot( m.01 )
```



useful model fit functions

```
coefficients( m.01 ) # model coefficients
```

```
## (Intercept)          x  
##   -4.460267    2.114140
```

```
confint( m.01, level=0.95) # CIs for model parameters
```

```
##              2.5 %   97.5 %  
## (Intercept) -12.087946 3.167412  
## x              1.983008 2.245272
```

```
# not run because of long output
```

```
# anova( m.01 ) # anova table
```

```
# fitted( m.01 ) # predicted values
```

```
# residuals( m.01 ) # residuals
```


```
# influence( m.01 ) # regression diagnostics
```

```
library( coefplot )
```

```
## Loading required package: ggplot2
```

```
m.02 <- lm( y ~ x1 + I(x1^2) + x2 + x3 )
```

```
coefplot(m.02)
```



pretty output

```
# install.packages( "memisc" )

library( memisc )

## Loading required package: lattice
##
## Attaching package: 'lattice'
## The following object is masked from 'package:boot':
##
##      melanoma
## Loading required package: MASS
##
## Attaching package: 'memisc'
## The following objects are masked from 'package:stats':
##
##      contr.sum, contr.treatment, contrasts
## The following object is masked from 'package:base':
##
##      as.array
m.02 <- lm( y ~ x + I(x^2) ) # quadratic term
m.03 <- lm( y ~ x - 1 )     # no intercept term

pretty.table <- mtable("Model 1"=m.01,"Model 2"=m.02,"Model 3"=m.03,
                      summary.stats=c("R-squared","F","p","N"))

pretty.table

##
## Calls:
## Model 1: lm(formula = y ~ x)
## Model 2: lm(formula = y ~ x + I(x^2))
## Model 3: lm(formula = y ~ x - 1)
##
## =====
##              Model 1      Model 2      Model 3
## -----
## (Intercept) -4.460      -3.122
##              (3.844)      (5.866)
## x            2.114***      2.035***      2.048***
##              (0.066)      (0.268)      (0.033)
## I(x^2)              0.001
##              (0.003)
## -----
## R-squared      0.9        0.9        1.0
## F              1023.6      507.1      3885.3
## p              0.0        0.0        0.0
## N              100        100        100
```

```
## =====
```

specification

```
summary( lm( y ~ x1 + x2 + x3 ) )
```

```
##
## Call:
## lm(formula = y ~ x1 + x2 + x3)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -46.466 -12.108  -0.452   11.528   53.562
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  -4.4359      3.9019  -1.137   0.258
## x1             2.1345      0.2175   9.812 3.77e-16 ***
## x2             0.2145      2.1115   0.102   0.919
## x3            -0.1826      1.8964  -0.096   0.923
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 19.27 on 96 degrees of freedom
## Multiple R-squared:  0.9126, Adjusted R-squared:  0.9099
## F-statistic: 334.3 on 3 and 96 DF,  p-value: < 2.2e-16
```

```
# add different functional forms
```

```
# square x1
```

```
summary( lm( y ~ x1 + x1^2 + x2 + x3 ) ) # not right
```

```
##
## Call:
## lm(formula = y ~ x1 + x1^2 + x2 + x3)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -46.466 -12.108  -0.452   11.528   53.562
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  -4.4359      3.9019  -1.137   0.258
## x1             2.1345      0.2175   9.812 3.77e-16 ***
## x2             0.2145      2.1115   0.102   0.919
## x3            -0.1826      1.8964  -0.096   0.923
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 19.27 on 96 degrees of freedom
## Multiple R-squared:  0.9126, Adjusted R-squared:  0.9099
## F-statistic: 334.3 on 3 and 96 DF,  p-value: < 2.2e-16
```



```
summary( lm( y ~ x1 + I(x1^2) + x2 + x3 ) ) # like this
```

```
##
## Call:
## lm(formula = y ~ x1 + I(x1^2) + x2 + x3)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -45.985 -12.295  -0.777   10.520   52.647
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -3.080226   5.968037  -0.516   0.607
## x1           2.058652   0.333357   6.176 1.62e-08 ***
## I(x1^2)       0.000785   0.002606   0.301   0.764
## x2           0.248586   2.124575   0.117   0.907
## x3          -0.153319   1.907933  -0.080   0.936
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 19.36 on 95 degrees of freedom
## Multiple R-squared:  0.9127, Adjusted R-squared:  0.9091
## F-statistic: 248.4 on 4 and 95 DF,  p-value: < 2.2e-16
```

```
summary( lm( y ~ log(x1) + x2 + x3 ) ) # log of x1 in formula works fine
```

```
##
## Call:
## lm(formula = y ~ log(x1) + x2 + x3)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -62.137 -18.258  -0.396   17.223   66.623
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -51.9951    12.2681  -4.238 5.18e-05 ***
## log(x1)       23.6072     4.8932   4.824 5.28e-06 ***
## x2          -13.3240     1.5229  -8.749 7.21e-14 ***
## x3           -0.6644     2.4076  -0.276   0.783
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 24.47 on 96 degrees of freedom
## Multiple R-squared:  0.8592, Adjusted R-squared:  0.8548
## F-statistic: 195.2 on 3 and 96 DF,  p-value: < 2.2e-16
```

```
# interactions
```

```
summary( lm( y ~ x1 + x2 ) )
```

```
##
## Call:
## lm(formula = y ~ x1 + x2)
##
```

```
## Residuals:
##      Min       1Q   Median       3Q      Max
## -46.759 -11.963  -0.287  11.304  53.458
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  -4.4241     3.8800  -1.140   0.257
## x1             2.1349     0.2164   9.866 2.61e-16 ***
## x2             0.2112     2.1004   0.101   0.920
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 19.17 on 97 degrees of freedom
## Multiple R-squared:  0.9126, Adjusted R-squared:  0.9108
## F-statistic: 506.6 on 2 and 97 DF,  p-value: < 2.2e-16
summary( lm( y ~ x1 + x2 + I(x1*x2) ) )
```

```
##
## Call:
## lm(formula = y ~ x1 + x2 + I(x1 * x2))
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -45.987 -11.869  -0.624  10.218  51.667
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -1.54690     5.89977  -0.262   0.794
## x1             2.05630     0.24852   8.274 7.45e-13 ***
## x2             1.06246     2.48179   0.428   0.670
## I(x1 * x2)  -0.01601     0.02468  -0.649   0.518
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 19.23 on 96 degrees of freedom
## Multiple R-squared:  0.913, Adjusted R-squared:  0.9103
## F-statistic: 335.9 on 3 and 96 DF,  p-value: < 2.2e-16
summary( lm( y ~ x1*x2 ) ) # shortcut
```

```
##
## Call:
## lm(formula = y ~ x1 * x2)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -45.987 -11.869  -0.624  10.218  51.667
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -1.54690     5.89977  -0.262   0.794
## x1             2.05630     0.24852   8.274 7.45e-13 ***
## x2             1.06246     2.48179   0.428   0.670
## x1:x2        -0.01601     0.02468  -0.649   0.518
```

```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 19.23 on 96 degrees of freedom
## Multiple R-squared:  0.913, Adjusted R-squared:  0.9103
## F-statistic: 335.9 on 3 and 96 DF,  p-value: < 2.2e-16

# dummy variables

summary( lm( y ~ x1 + x2 + x3 + dum ) ) # drop one level

##
## Call:
## lm(formula = y ~ x1 + x2 + x3 + dum)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -42.163 -12.245  -0.965  10.027  51.709
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  -2.6334     4.7769  -0.551   0.583
## x1             2.1338     0.2191   9.741 7.24e-16 ***
## x2             0.1342     2.1294   0.063   0.950
## x3            -0.2050     1.9431  -0.105   0.916
## dumNJ         -0.7546     5.3030  -0.142   0.887
## dumNY         -2.3549     5.4269  -0.434   0.665
## dumPA         -6.6512     5.4422  -1.222   0.225
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 19.4 on 93 degrees of freedom
## Multiple R-squared:  0.9142, Adjusted R-squared:  0.9087
## F-statistic: 165.1 on 6 and 93 DF,  p-value: < 2.2e-16

summary( lm( y ~ x1 + x2 + x3 + dum - 1 ) ) # keep all, drop intercept

##
## Call:
## lm(formula = y ~ x1 + x2 + x3 + dum - 1)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -42.163 -12.245  -0.965  10.027  51.709
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## x1             2.1338     0.2191   9.741 7.24e-16 ***
## x2             0.1342     2.1294   0.063   0.950
## x3            -0.2050     1.9431  -0.105   0.916
## dumMA         -2.6334     4.7769  -0.551   0.583
## dumNJ         -3.3880     5.3696  -0.631   0.530
## dumNY         -4.9882     5.1799  -0.963   0.338
## dumPA         -9.2846     5.6346  -1.648   0.103
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 19.4 on 93 degrees of freedom
## Multiple R-squared:  0.9759, Adjusted R-squared:  0.9741
## F-statistic: 538.6 on 7 and 93 DF,  p-value: < 2.2e-16
```

standardized regression coefficients (beta)

```
# install.packages( "lm.beta" )
```

```
library( lm.beta )
```

```
m.01.beta <- lm.beta( m.01 )
```

```
summary( m.01.beta )
```

```
##
## Call:
## lm(formula = y ~ x)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -47.000 -11.932  -0.287   11.197   53.447
##
## Coefficients:
##              Estimate Standardized Std. Error t value Pr(>|t|)
## (Intercept) -4.46027      0.00000      3.84369   -1.16   0.249
## x              2.11414      0.95531      0.06608   31.99 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 19.07 on 98 degrees of freedom
## Multiple R-squared:  0.9126, Adjusted R-squared:  0.9117
## F-statistic: 1024 on 1 and 98 DF,  p-value: < 2.2e-16
```

```
# coef( m.01.beta )
```

```
# note the standard error is not standardized - describes regular coefficients not standardized
```

```
summary( m.01 )
```

```
##
## Call:
## lm(formula = y ~ x)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -47.000 -11.932  -0.287   11.197   53.447
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -4.46027      3.84369   -1.16   0.249
## x              2.11414      0.06608   31.99 <2e-16 ***
```

```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 19.07 on 98 degrees of freedom
## Multiple R-squared:  0.9126, Adjusted R-squared:  0.9117
## F-statistic: 1024 on 1 and 98 DF,  p-value: < 2.2e-16
```

or just use the formula:

```
lm.beta <- function( my.mod )
{
  b <- summary(my.mod)$coef[-1, 1]
  sx <- sd( my.mod$model[, -1] )
  sy <- sd( my.mod$model[, 1] )
  beta <- b * sx/sy
  return(beta)
}
```

```
coefficients( m.01 )
```

```
## (Intercept)          x
## -4.460267    2.114140
```

```
lm.beta( m.01 )
```

```
## [1] 0.9553146
```

robust standard errors

```
# install.packages( "sandwich" )
# install.packages( "lmtest" )
```

```
library(sandwich)
library(lmtest)
```

```
## Loading required package: zoo
```

```
##
```

```
## Attaching package: 'zoo'
```

```
## The following objects are masked from 'package:base':
```

```
##
```

```
##      as.Date, as.Date.numeric
```

```
m.01 <- lm( y ~ x )
```

```
# REGULAR STANDARD ERRORS
```

```
summary( m.01 ) # not-robust
```

```
##
```

```
## Call:
```

```
## lm(formula = y ~ x)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -47.000 -11.932  -0.287   11.197   53.447
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -4.46027     3.84369   -1.16   0.249
## x             2.11414     0.06608   31.99 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 19.07 on 98 degrees of freedom
## Multiple R-squared:  0.9126, Adjusted R-squared:  0.9117
## F-statistic: 1024 on 1 and 98 DF,  p-value: < 2.2e-16

# ROBUST STANDARD ERRORS

# reproduce the Stata default
coeftest( m.01, vcov=vcovHC(m.01,"HC1") )      # robust; HC1 (Stata default)

##
## t test of coefficients:
##
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -4.460267     3.945755  -1.1304   0.2611
## x             2.114140     0.065438  32.3076   <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

# ROBUST STANDARD ERRORS

# check that "sandwich" returns HCO
coeftest(m.01, vcov = sandwich)                # robust; sandwich

##
## t test of coefficients:
##
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -4.46027     3.90610  -1.1419   0.2563
## x             2.11414     0.06478  32.6356   <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

coeftest(m.01, vcov = vcovHC(m.01, "HC0"))      # robust; HCO

##
## t test of coefficients:
##
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -4.46027     3.90610  -1.1419   0.2563
## x             2.11414     0.06478  32.6356   <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
# ROBUST STANDARD ERRORS
```

```
# check that the default robust var-cov matrix is HC3
```

```
coeftest(m.01, vcov = vcovHC(m.01)) # robust; HC3
```

```
##
```

```
## t test of coefficients:
```

```
##
```

```
##           Estimate Std. Error t value Pr(>|t|)
```

```
## (Intercept) -4.460267   4.017243 -1.1103   0.2696
```

```
## x           2.114140   0.066632 31.7288   <2e-16 ***
```

```
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
coeftest(m.01, vcov = vcovHC(m.01, "HC3")) # robust; HC3 (default)
```

```
##
```

```
## t test of coefficients:
```

```
##
```

```
##           Estimate Std. Error t value Pr(>|t|)
```

```
## (Intercept) -4.460267   4.017243 -1.1103   0.2696
```

```
## x           2.114140   0.066632 31.7288   <2e-16 ***
```

```
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```